

EZPWD Reed-Solomon Version 1.2.0

Perry Kundert

March 25, 2015

Contents

| | | |
|----------|---|-----------|
| 1 | Reed-Solomon Loss/Error Correction | 2 |
| 1.1 | Performance | 2 |
| 1.1.1 | 25-75% Faster Than Competing Codecs | 2 |
| 1.2 | Availability | 3 |
| 1.2.1 | C++ Build Requires Header Source Only | 3 |
| 1.2.2 | Javascript Hosted by MaxCDN | 3 |
| 1.2.3 | Python Interface using Swig | 3 |
| 1.3 | Licensing | 4 |
| 1.3.1 | GPLv3 Licensing | 4 |
| 1.3.2 | Commercial Licensing and Pricing | 4 |
| 1.4 | Enhancements | 5 |
| 1.4.1 | Rejects impossible error position | 5 |
| 1.4.2 | Shared data tables w/ optional locking | 6 |
| 1.5 | ezpwd::RS<...>: C++ Reed-Solomon API | 6 |
| 2 | RSKEY: Data Key API | 6 |
| 2.1 | Javascript Library: js/ezpwd/rskey.js | 7 |
| 2.2 | RSKEY Demo: http://rskey.hardconsulting.com | 8 |
| 2.3 | Example Node.JS: Encrypted Gift Card Codes | 8 |
| 2.3.1 | Client Website RSKEY Implementation | 9 |
| 2.3.2 | Server Node.js Encryption Implementation | 11 |
| 3 | EZCOD: Location Code API | 13 |
| 3.1 | Javascript Library: js/ezpwd/ezcod.js: | 14 |
| 3.2 | Python Library: ezpwd_reed_solomon | 15 |
| 3.2.1 | ezpwd_reed_solomon.ezcod | 15 |
| 3.2.2 | Asynchronous Loading | 18 |
| 3.3 | Robustness | 19 |

| | | |
|----------|---|-----------|
| 3.4 | Precision | 19 |
| 3.5 | EZCOD Demo: http://ezcod.com | 20 |
| 4 | RSPWD: Password Correction API | 20 |
| 4.1 | Javascript Library: js/ezpwd/rspwd.js | 20 |

1 Reed-Solomon Loss/Error Correction

Error and erasure detection and correction for C++, Javascript and Python programs. Based on Phil Karn's excellent implementation (as used by the Linux kernel), converted to C++.

Several C++ and Javascript utilities implemented using Reed-Solomon are provided.

The Javascript APIs are produced from the C++ implementation using `emscripten`, and are very performant. All Javascript is available for inclusion in websites via MaxCDN.

1.1 Performance

The Reed-Solomon codec's performance is often a critical determinant in product performance, and often influences expensive decisions such as CPU selection for embedded applications.

The `ezpwd::RS<...>` codec corrects data in-place to avoid unnecessary copying, using shared data tables between RS instances with compatible size, parity and Galois field parameters. The use of C++ template parameters for size-related Reed Solomon parameters enables the compiler to emit optimal code for all internal table lookups, sometimes dramatically improving performance.

1.1.1 25-75% Faster Than Competing Codecs

EZPWD Reed Solomon performs R-S correction operations (using g++ on an Intel i7) about 75% faster than Phil Karn's general case code, and about 25% faster than Phil's optimized code for 8-bit/CCSDS symbols. It also tests at about 25% faster than the Shifra C++ Reed-Solomon library. It appears that, for certain applications at least, EZPWD Reed-Solomon may be the fastest Open Source general-purpose C++ Reed-Solomon library presently available.

1.2 Availability

Source code is hosted at <https://github.com/pjkundert/ezpwd-reed-solomon>. It is heavily tested under both g++ and clang++, with full optimization enabled.

1.2.1 C++ Build Requires Header Source Only

Carefully implemented as C++11 standard inline code, to require no C++ object-code or library build procedure. As a result, integration into existing C++ code-bases and build systems is extremely simple.

1.2.2 Javascript Hosted by MaxCDN

Production minified Javascript is served by MaxCDN, using a URL like: <https://cdn.rawgit.com/pjkundert/ezpwd-reed-solomon/v1.2.0/js/ezpwd/ezcod.js>.

For example, to use the EZCOD position encoding API in your website, add this at the end of your website template, before your application.js:

```
<script
  src="//cdn.rawgit.com/pjkundert/ezpwd-reed-solomon/v1.2.0/js/ezpwd/ezcod.js">
</script>
```

Note the current <VERSION> number encoded in the URI as `=.../v<VERSION>/...`.

Your application may permanently access any current or historical version of the EZPWD Reed-Solomon Javascript which has a Tag in the official Git repo, by encoding the <VERSION> number in the URL:

```
//cdn.rawgit.com/.../v<VERSION>/...
```

1.2.3 Python Interface using Swig

A high-performance Python 2/3 interface is provided, using Swig (requires at least Swig Version 3.0.5 for C++11 support). Since the Python module is generated from the C++ interface, it must be generated and compiled using the appropriate target OS, Python and C++ compiler implementation. Therefore, there is no `pip install ezpwd_reed_solomon` available.

To build the `ezpwd_reed_solomon` package, obtain the source from <https://github.com/pjkundert/ezpwd-reed-solomon>, and build with the version of Python you wish to support:

```
$ cd ezpwd-reed-solomon/swig/python
$ python setup.py install
```

Presently, only `ezpwd_reed_solomon.ezcod` is available; see the EZCOD: Location Code API section for Python API details.

- Swig 3.0.5+

To install the Python API, you'll need a modern Swig.

On Mac, use homebrew:

```
$ brew install swig
```

On Linux Debian or Ubuntu Linux systems, you should be able to use something like this (other Linux variants should have similar package installation facilities):

```
$ apt-get -u install autoconf autogen libpcrc3-dev bison yodl
$ git clone git@github.com:swig/swig.git
$ cd swig
$ autoconf
$ ./autogen.sh
$ ./configure && make && make install
```

1.3 Licensing

All `ezpwd-reed-solomon` code is available under both GPLv3 and Commercial licenses. Phil's original Reed-Solomon code is LGPL, so my Reed-Solomon implementation in `ezpwd-reed-solomon/c++/ezpwd/rs` (which uses Phil's, with some improvements and conversion to C++) is available under the terms of the LGPL.

1.3.1 GPLv3 Licensing

If your application complies with the terms of the GPLv3, then you can use EZPWD Reed-Solomon based APIs without cost. All users of your software (eg. an installed application) or "software as a service" (eg. a website) must have access to all of the software source code so they can freely modify, rebuild and run the software. Any modifications to underlying GPLv3 software (ie. EZPWD) must also be made available.

1.3.2 Commercial Licensing and Pricing

If you use any of the EZPWD Reed-Solomon based APIs in your product but you don't wish to make your product's or website's source code available,

then a Commercial license from Hard Consulting Corporation (hardconsulting.com) is available. The pricing breakdown is as follows (in CAD\$):

| Users avg. (monthly) | Price CAD\$ | Support CAD\$/yr | Included application assistance |
|-------------------------|----------------|---------------------|---------------------------------|
| <1K | 0 | 25 | Interesting project? ask... :) |
| 1K-10K | 100 | 25 | 1 hour |
| 10K-1M | 1000 | 250 | 4 hours |
| >1M | 10000 | 2500 | 16 hours |

Use of the EZCOD robust geolocation encoding module of EZPWD Reed-Solomon is free, forever, for any application. It is available under both GPLv3 and free Commercial licenses, and may even be re-implemented freely in any language, so long as it remains compatible (includes the Reed-Solomon error correction, and equivalent encoding and decoding of Latitude and Longitude coordinates).

Call us at +1-780-970-8148 or email us at info@hardconsulting.com to discuss your application.

1.4 Enhancements

Several enhancements have been made to Phil's implementation.

1.4.1 Rejects impossible error position

Phil's version allows the R-S decode to compute and return error positions with the unused portion of the Reed-Solomon codeword. We reject these solutions, as they provide indication of a failure.

The supplied data and parity may not employ the full potential codeword size for a given Reed-Solomon codec. For example, and RS(31,29) codec is able to decode a codeword of 5-bit symbols containing up to 31 data and parity symbols; in this case, 2 parity symbols ($31-29 == 2$).

If we supply (say) 9 data symbols and 2 parity symbols, the remaining 20 symbols of unused capacity are effectively filled with zeros for the Reed-Solomon encode and decode operations.

If we decode such a codeword, and the R-S Galois field solution indicates an error positioned in the first 20 symbols of the codeword (an impossible situation), we reject the codeword and return an error.

1.4.2 Shared data tables w/ optional locking

Instead of re-computing all of the required data tables used by the Reed-Solomon computations, every instance of `RS<CAPACITY,*>` with compatible Galois polynomial parameters shares a common set of tables. Furthermore, every instance of `RS<CAPACITY,PAYLOAD>` w/ compatible Galias polynomial parameters shares the tables specific to the computed number of parity symbols.

The initialization of these tables is protected by a Mutex primitive and Guard object. These default to 'int' (NO-OP), but if a threading mutex and guard are provided, the shared initialization is thread-safe.

1.5 `ezpwd::RS<...>`: C++ Reed-Solomon API

C++ implementation of Reed-Solomon codec. Fully implemented as inline code, in C++ header files. Highly performant, in both C++ and Javascript.

```
#include <ezpwd/rs>

// Reed Solomon codec w/ 255 symbols, up to 251 data, 4 parity symbols
ezpwd::RS<255,251> rs;

std::vector<uint8_t> data;

// ... fill data with up to 251 bytes ...

rs.encode( data ); // Add 4 Reed-Solomon parity symbols (255-251 == 4)

// ... later, after data is possibly corrupted ...

int fixed = rs.decode( data ); // Correct errors, discard 4 R-S parity symbols
```

2 RSKEY: Data Key API

Asking a user to reliably enter even a few bytes of data (eg. a product key or a redemption code) is, well, basically impossible. It is not reasonable to expect that someone will be able to perfectly read and enter a blob of random letters and numbers with 100% reliability.

Import `js/ezpwd/rskey.js` to use RSKEY error-corrected binary data input tokens in your application. Raw binary data (in Javascript string or

ArrayBuffer) can be encoded into an RSKEY for later entry by a user. Using built-in parity (extra validation) symbols, any errors or missing symbols can be detected and possibly recovered. An RSKEY that validates as correct can be trusted with a high degree of certainty, proportional to the number of excess parity symbols remaining (beyond those consumed by error detection and correction).

2.1 Javascript Library: `js/ezpwd/rskey.js`

```
rskey_<PARITY>_encode( <bytes>, data, [ sep ] ) -- encode data to RSKEY
rskey_<PARITY>_decode( <bytes>, key )           -- decode RSKEY
```

PARITY of 2-5 is supported, with a maximum capacity of 31-PARITY bytes of base-32 encoded data (raw data expands by the factor $(\text{<bytes>} * 8 + 4) / 5$ when base-32 encoded). With PARITY 2, the maximum capacity is 18 bytes; with PARITY 5, 16 bytes.

The `data` may be an ArrayBuffer of byte-length $\leq \text{<bytes>}$. If a string is supplied, it may be a hex string beginning with ‘0x...’ (all subsequent pairs of hex digits are used; any data beyond that is ignored). Otherwise, the string is decoded as utf-8 (of course, this means that you can’t supply a utf-8 string that starts with ‘0x’...).

The optional `sep` parameter (default 5) is the cluster size to separate the RSKEY into; 0 specifies no separators.

Load the `rskey.js` Javascript into your project:

```
<script
  src="//cdn.rawgit.com/pjkundert/ezpwd-reed-solomon/v1.2.0/js/ezpwd/rskey.js">
</script>
```

Use `rskey.js`’s API to encode your data into an easily human readable key. Call the `rskey_<PARITY>_encode` API (with PARITY 2-5), specify the number of bytes of data to encode in the RSKEY’s payload, and provide some data to encode (as a hex string “0x3344...”, or as a utf-8 string):

```
> rskey_5_encode( 12, "Mag.1ck" );
"9MGNE-BHHCD-MVY00-00000-MVRFN"
```

Later, you can decode it – even if the user adds an error or two (the ‘X’, below), or skips a few symbols (if some were unreadable, as indicated by an ‘_’, or the last few are simply not yet entered). Each error consumes 2 parity symbols, each erasure or missing symbol uses 1, therefore 1 error + 2 erasures results in 20% of parity remaining for validation:

```

> rskey_5_decode( 12, "9MGNE-BHHCD-MVY00-00000-MVRFN" )
{confidence: 100, data: ArrayBuffer, utf8: "Mag.1ck", hex: "0x4D61672E31636BCF80000000"}
> rskey_5_decode( 12, "9MGNE-BHHCD-MVY00-00X00-MVR" ) // 1 error, 2 not yet entered
{confidence: 20, data: ArrayBuffer, utf8: "Mag.1ck", hex: "0x4D61672E31636BCF80000000"}
> rskey_5_decode( 12, "9_GNE-BHH_D-MVY00-00X00-MVRFN" ) // 1 error, 2 unreadable w/ '_'
{confidence: 20, data: ArrayBuffer, utf8: "Mag.1ck", hex: "0x4D61672E31636BCF80000000"}

```

If you have raw numeric data (eg. record IDs, data HMACs, etc), use the `ArrayBuffer` interface. You can supply any type of raw data, up to the capacity of the `RSKEY` (12 bytes, in this case). Then, even if errors are introduced on entry, they will be recovered if the parity is sufficient, and the returned Object's `.data` property will be an `ArrayBuffer` containing the original binary data, which you can use a `TypedArray` to access:

```

> ia = new Int32Array([0x31323334, 0x41424344, 0x51525354])
[825373492, 1094861636, 1364349780]
> rskey_5_encode( 12, ia.buffer ) // raw capacity is 12 bytes, w/ 5 parity
"6GRK4-CA48D-142M2-KA98G-V2MYP"
> dec=rskey_5_decode( 12, "6GRK4-CA48D-142M2-KA98G-V2XXP" ) // XX are errors
{confidence: 20, data: ArrayBuffer, utf8: "4321DCBATSRQ", hex: "0x34333231444342415453"}
> new Int32Array( dec.data ) // recover original data
[825373492, 1094861636, 1364349780]

```

2.2 RSKEY Demo: <http://rskey.hardconsulting.com>

Try changing the Parity, Data Size and Data. Try changing the Key by entering some `_` (indicating a missing/invalid symbol). These are called Erasures in Reed-Solomon terms, and we can recover one Erasure with each Parity symbol. Try changing some Key values to incorrect values. These Reed-Solomon Errors each require 2 Parity symbols to detect and correct.

You can also access the Console (right click, select Inspect Element, click on "Console"), and enter the above `rskey_...` API example code.

2.3 Example Node.JS: Encrypted Gift Card Codes

Lets say you have an online Widget business, and generate gift cards. You average about 5000 unique visitors/month over the year, with a peak of 25000 around Christmas. You want to make your gift card redemption more reliable and secure, and less painful for your clients.

Your `RSKEY` license cost would be \$100, plus a \$25/yr support subscription, and you would have access to an hour of time with a support

developer to help you apply the `js/ezpwd/rskey.js` API to your website's gift card generation and redemption pages.

You decide to associate each gift card with the buyer's account (so you and the gift-card giver can know when the card is redeemed). So, each gift card RSKEY needs to contain:

- a 32-bit customer ID
- a 32-bit gift card ID

Using an RSKEY encoding 8 bytes of data, with 3 parity symbols, we get protection against 1 error or 2 erased symbols, with 1 parity symbol left over for validation.

See `rskey_node.js` for sample code (the communication of the JSON request and reply between the client Website and the Node.JS server is left as an exercise to the reader.)

2.3.1 Client Website RSKEY Implementation

On the client website, you would use something like:

```
<script
  src="//cdn.rawgit.com/pjkundert/ezpwd-reed-solomon/v1.2.0/js/ezpwd/rskey.js">
</script>
<script>
var client = {
  //
  // card_key_encode( card ) -- encrypt card's IDs on the server, return RSKEY
  // card_key_decode( key ) -- recover RSKEY, decrypt IDs on server, return card
  //
  //      These are run in the browser, and expect to call server methods that
  // run under Node.js back on the server.  For this demo, we'll all just run
  // here in Node.js...
  //
  card_key_encode: function( card ) {
    // Get the server to encrypt the card IDs
    server.card_keydata_encode( card );
    // Produce the RSKEY from the card's keydata w/ Uint8Array's ArrayBuffer
    card.key = rskey_3_encode( 8, new Uint8Array( card.keydata ).buffer, 4 );
    return card.key;
  },
}
```

```

card_key_decode: function( key ) {
    // Decode the ASCII key; will raise an Exception if decode fails
    var keyinfo = rskey_3_decode( 8, key );

    // Convert ArrayBuffer (as Uint8Array) to plain javascript Array
    var keyuint8 = new Uint8Array( keyinfo.data );
    var keydata = Array( 8 );
    for ( var i = 0; i < 8; ++i )
        keydata[i] = keyuint8[i];

    // Get the server to decrypt the card.keydata, return the card IDs
    return server.card_keydata_decode({ keydata: keydata });
}
}

// Your first customer ever, buys his first gift card!
card = {
    id: 0,
    customer: { id: 0 },
}

// Encode the card IDs to RSKEY
card_key = client.card_key_encode( card );
// ==> {
//   customer: { id: 0 },
//   id: 0,
//   keydata: [ 185, 124, 29, 95, 168, 45, 159, 113 ],
//   key: 'P5X1-TPW8-5NFP-2M7G'
// }
//
// "P5X1-TPW8-5NFP-2M7G" is printed/emailed on gift card
//

```

Later on, the gift card recipient comes back to the website and enters the gift-card key during checkout, mistyping some symbols, and using lower-case and alternative whitespace (the base-32 encoding handles the Z/z/2, S/s/5, I/i/1 and O/o/0 substitutions (these symbols are equivalent in EZPWD base-32); the W/v substitution is an error):

```
// Decode the customer-entered data using the same RSKEY parameters:
```

```

//                      error:          v
//                      equivalents:    v v          v v
//                      original: "P5X1-TPW8-5NFP-2M7G"
card_dec = client.card_key_decode( "psxi tpv8 snfp zm7g" );
// ==> {
//   keydata: [ 185, 124, 29, 95, 168, 45, 159, 113 ],
//   customer: { id: 0 },
//   id: 0
// }
//
// This is gift card ID 0, purchased by our very first customer ID 0! Find out
// what that gift card is still worth, and apply it to the order...
//

```

2.3.2 Server Node.js Encryption Implementation

All encryption should take place on the server, with a secret symmetric encryption key (which should not be stored in the repo! Use other secure key storage, or existing key material already on the server). Encrypt on the server using an appropriate cipher that either encrypts all 64 bits as a block (such as `blowfish`).

```

/*
 * rskey_node.js -- Demonstrate use of rskey in Node.js application
 *
 *   Node.js "crypto" uses the Buffer type to manipulate binary data. The
 *   rskey library uses ArrayBuffer, because it is intended to be used in both
 *   Node.js and Browser Javascript applications.
 *
 *   The server will expect an Object containing (at least) card.id and
 *   card.customer.id, and produce/consume card.keydata.
 */
var crypto          = require( "crypto" );
var crypto_algo     = 'blowfish'; // 64-bit block cipher
var crypto_secret   = 'not.here'; // Super secret master key; don't keep in Git...

var server = {
  //

```

```

// card_keydata_encode -- Encipher card IDs into card.keydata Array
// card_keydata_decode -- Decipher card IDs from card.keydata Array
//
//      Run these on your server (of course, keeping crypto_secret... secret.)
//
card_keydata_encode: function( card ) {
    // Create Buffer containing raw card ID data
    var buf          = new Buffer( 8 );
    buf.writeUInt32LE( card.customer.id,    0 );
    buf.writeUInt32LE( card.id,             4 );

    // Encrypt the Buffer of keydata
    var encipher     = crypto.createCipher( crypto_algo, crypto_secret );
    encipher.setAutoPadding( false ); // must use exact 64-bit blocks
    var enc          = Buffer.concat([
        encipher.update( buf ),
        encipher.final()
    ]);

    // Return card w/ encrypted IDs as plain Javascript Array in .keydata
    card.keydata     = enc.toJSON().data; // {type: 'Buffer', data: [1,2,...]}
    return card;
},

card_keydata_decode: function( card ) {
    if ( card.keydata.length != 8 )
        throw "Expected 8 bytes of card.keydata, got: " + card.keydata.length;

    // Decrypt the Buffer of keydata
    var decipher     = crypto.createDecipher( crypto_algo, crypto_secret );
    decipher.setAutoPadding( false ); // must use exact 64-bit blocks
    var dec          = Buffer.concat([
        decipher.update( new Buffer( card.keydata ) ),
        decipher.final()
    ]);

    // Recover raw card IDs from Buffer
    if ( card.customer == undefined )
        card.customer = {};
    card.customer.id= dec.readUInt32LE( 0 );

```

```

        card.id          = dec.readUInt32LE( 4 );
        return card;
    }
};

```

Assuming that an attacker does not have access to the encryption key used by the server to encrypt the customer and card IDs in a single 64-bit block, then the probability of a fake key being produced and accepted is vanishingly small.

Lets assume that they **do** know that you are using EZPWD Reed-Solomon, and therefore always present RSKEYs that are valid R-S code-words. Furthermore, lets assume that you have alot of customers (> 2 billion), so your 32-bit customer ID is likely to accidentally match a valid customer with a probability $> 50\%$.

The decrypted customer and card IDs must be correct – match a valid customer and card ID. Since it is unlikely for each customer to generate more than a handful of gift cards, the probability that the 32-bit card ID will accidentally decrypt to any given value is $1/2^{32}$ (1 in ~ 4 billion). The combined 64-bit RSKEY (remember: all data must be encrypted with a block cipher) indexes a sparsely populated array of valid values; given a number in the range $(0, 2^{64}]$, only every 4-billionth value will turn out to be valid (much less than that, in realistic scenarios).

Therefore, an attacker must generate and try more than 2 billion valid RSKEYs before they have a 50% chance of stumbling upon one that matches a valid gift card, given the above (generous) assumptions. Even if you don't rate-limit your card redemption API, you might notice that your server is saturated with gift-card redemption requests. Assuming that your server can process 1000 redemptions per second, it would take the attacker 23 days (2,000,000 seconds) to have a 50% chance of finding his first valid fake key. So, I recommend rate-limiting your gift-card redemption API to 10 request per second, increasing the time to 6 years.

Therefore, using RSKEY and a simple encoding scheme presents an effective, robust and secure means of generating and redeeming gift-card codes.

Customer aggravation due to mis-typed codes is reduced, increasing the likelihood of return visits and positive reviews.

3 EZCOD: Location Code API

To specify the location of something on the surface of the earth, a Latitude, Longitude pair is typically used. To get within $\pm 3m$, a Latitude, Longitude

pair with at least 5 digits of precision after the decimal point is required.

So, to specify where my daughter Amarissa was born, I can write down the coordinate:

```
53.655832,-113.625433
```

This is both longer and more error prone than writing the equivalent EZCOD:

```
R3U 1JU QUY.0
```

If a digit is wrong in the Latitude or Longitude coordinate, the amount of error introduced is anywhere from a few centimeters to many kilometers:

```
53.655832,-113.62543X == centimeters error
53.655832,-1X3.625433 == many kilometers error
```

EZCOD uses error/erasure correction to correct for up to 1 known missing (erased) symbol by default, with greater erasure/error detection and correction optionally available.

3.1 Javascript Library: `js/ezpwd/ezcod.js`:

```
ezcod_3_10_encode( lat, lon, [ symbols ] ) -- encode location to EZCOD
ezcod_3_10_decode( ezcod )                -- decode EZCOD to position
```

There are three variants provided:

- `ezcod_3_10_...` – 1 parity symbol
- `ezcod_3_11_...` – 2 parity symbols
- `ezcod_3_12_...` – 3 parity symbols

Load the `ezcod.js` Javascript into your project:

```
<script
  src="//cdn.rawgit.com/pjkundert/ezpwd-reed-solomon/v1.2.0/js/ezpwd/ezcod.js">
</script>
```

To encode a position of center of the Taj Mahal dome to 3m accuracy (9 position symbols, the default) and 20mm accuracy (12 symbols), and with 3 parity symbols (5-nines confidence):

```
> ezcod_3_12_encode( 27.175036985, 78.042124565 ) // default: 3m (9 symbols)
"MMF BBF GC1.2U2"
> ezcod_3_12_encode( 27.175036985, 78.042124565, 12 ) // 20mm (12 symbols)
"MMF BBF GC1 A16.1VD"
```

Later, if the EZCOD is entered, errors and erasures are transparently corrected, up to the capacity of the Reed-Solomon encoded parity:

```
> ezcod_3_12_decode( "MMF BBF GC1 A16.1VD" )
Object {confidence: 100, latitude: 27.17503683641553, longitude: 78.04212455637753,
        accuracy: 0.020401379521588606}
> ezcod_3_12_decode( "MMF BBF GC1 A16.1" ) // missing some parity
Object {confidence: 34, latitude: 27.17503683641553, longitude: 78.04212455637753,
        accuracy: 0.020401379521588606}
> ezcod_3_12_decode( "mmf-bbf-Xc1-a16.1vd" ) // An error
Object {confidence: 34, latitude: 27.17503683641553, longitude: 78.04212455637753,
        accuracy: 0.020401379521588606}
```

Try it at ezcod.com. Switch to “EZCOD 3:12”, and enter “mmf-bbf-Xc1-a16.1vd” as the EZCOD. You will see a computed accuracy of 20.4mm, and observe that the ‘X’ (error) is corrected to “G”. (The website defaults to 9 digits of precision, so it will re-encode the position, discarding the extra precision.)

3.2 Python Library: `ezpwd_reed_solomon`

The Python `ezpwd_reed_solomon` package contains an `ezcod` sub-module:

3.2.1 `ezpwd_reed_solomon.ezcod`

Classes are provided to produce three variants of EZCOD by default: 3m (9 symbols) of location accuracy, plus 1, 2 or 3 Reed-Solomon parity symbols. They are named `ezcod_3_10`, `ezcod_3_11` and `ezcod_3_12`, respectively, indicating the default 3m accuracy, and the total number of symbols.

```
$ python
>>> from ezpwd_reed_solomon import ezcod
```

The API supports the following methods and attributes:

- `ezcod_3_{10,11,12}("<EZCOD>"|[lat],[lon],[seper],[chunk]]])`
Creates an `<ezcod>` instance containing the specified geolocation (defaults to latitude 0.0, longitude 0.0, '.' separator and chunk 3). If a string is supplied, it is decoded (if possible; an Exception is raised if the provided EZCOD is invalid).

```
>>> EZCOD = ezcod.ezcod_3_12( 53.5, -113.8 )
>>> print EZCOD
R3U 06B MJ3.JXR (100%) == +53.5000000000, -113.8000000000 +/- 0.00mm
```

- `ezcod_3_{10,11,12}.encode([precision])`
Encodes the current `ezcod310,11,12`'s `.latitude` and `.longitude` to the given number of symbols of precision (default: 9, or 3m). The accuracy may be anywhere from 1 to 12 (20mm accuracy) symbols.

```
>>> print EZCOD.encode( 12 )
R3U 06B MJ3 EDD.K56
```

- `ezcod_3_{10,11,12}.decode("<EZCOD>")`
Any variant of `ezcod310,11,12` can decode a valid EZCOD, so long as it contains a '.' or '!' to separate the position and R-S parity symbols. The percentage certainty is returned – the proportion of R-S parity symbols that remain unused after error detection and correction. A value of 0 indicates that the EZCOD's R-S decoding did not fail, but no parity symbols remain in excess to verify its validity.

```
>>> print EZCOD.decode( "r3u 06b mj3 edd.k56" )
100
>>> EZCOD.latitude
53.49999999627471
>>> print EZCOD.decode( "r3u 06b m_3 edd.k56" )
67
>>> print EZCOD.decode( "r3u 06b mX3 edd.k56" )
34
>>> print EZCOD
R3U 06B MJ3.JXR ( 34%) == +53.4999999963, -113.8000000734 +/- 19.4mm
>>>
```

If any symbols are unknown, replace them with either `_` or `?` to indicate that they are erasures (and consume only a single symbol of R-S parity

to correct). Any undetected erroneous symbol corrected by the R-S codec consumes 2 parity symbols. A failure to decode (too many errors or erasures) will raise a `RuntimeError` exception:

```
>>> EZCOD.decode( "r3u 06b mj3 __d.__6" )
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  ...
RuntimeError: ezpwd::ezcod::decode: Error correction failed; too many erasures
>>> EZCOD.decode( "r3u 06b mj3 eXd.__6" )
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
  ...
RuntimeError: ezpwd::ezcod::decode: Error correction failed; R-S decode failed
```

- `ezcod_3_{10,11,12}` Instance Attributes

The following attributes are available in each `ezcod_3_{10,11,12}` instance:

| Attribute | Type | Range | Description |
|----------------------------|-------|-------------|---|
| latitude | float | [-90,90] | Geographical position in degrees |
| longitude | float | [-180,180] | “ |
| latitude _{error} | float | [0,inf] | Axis error in meters |
| longitude _{error} | float | [0,inf] | “ |
| accuracy | float | [0,inf] | Average of error ellipse axes in meters |
| confidence | int | [0,100] | Percentage of parity symbols remaining |
| chunk | int | [0,6] | Spaces every ‘chunk’ position symbols |
| seper | int | [0,3] | SEP _{NONE} , SEP _{DOT} , SEP _{BANG} , SEP _{SPACE} |
| SEP_NONE | int | 0 | Output no position/parity separator |
| SEP_DOT | int | 1 (default) | Output a ‘.’ position/parity separator |
| SEP_BANG | int | 2 | Output a ‘!’ position/parity separator |
| SEP_SPACE | int | 3 | Output a ‘ ‘ position/parity separator |

It is recommended to use either `SEP_DOT` (default) or `SEP_BANG`, so that the EZCOD parser can unambiguously determine the total EZCOD size, and the number of parity symbols to expect.

3.2.2 Asynchronous Loading

Emscripten-generated code must have its run-time initialized before it can be called. If you get Javascript resources normally, they will load asynchronously, but be run in the order you load them so the Emscripten run-time will be safely initialized before your application's Javascript runs.

If you load other Javascript libraries like jQuery and your application.js, and you load `ezcod.js` asynchronously, you must ensure that they do not use any Emscripten libraries (such as `ezcod.js`) until their run-time initialization is complete. Our Emscripten-based libraries are completely self-contained, so you can use the `<script onload...>` to signal jQuery to trigger its `on('ready', ...)` event. Regardless of whether `jquery.min.js` or `ezcod.js` loads first, this code will ensure that your `app.js on('ready', ...)` event will not fire until `ezcod.js` has its Emscripten run-time initialized:

```
<script type="text/javascript">
  // Bindings for Emscripten initialization detection.
  var jquery_release = function() {
    console.log( "Emscripten run-time initialized before jQuery loaded" );
    jquery_loaded = function() {}; // nothing left to do after jquery loads
  };
  var jquery_loaded = function() {
    console.log( "Emscripten run-time initialize blocking jQuery..." );
    $.holdReady(true);           // force delay of jQuery.on( 'ready', ...
    jquery_release = function() {
      console.log( "Emscripten run-time initialized; jQuery released" );
      $.holdReady(false);       // ... 'til Emscripten runtime initialized
    };
  };
</script>
<script async onload="jquery_release()"
  src="//cdn.rawgit.com/pjkundert/ezpwd-reed-solomon/v1.2.0/js/ezpwd/ezcod.js">
</script>
<script defer onload="jquery_loaded()"
  src="//ajax.googleapis.com/ajax/libs/jquery/2.1.3/jquery.min.js">
</script>
<script defer
  src="js/app.js">
</script>
```

3.3 Robustness

All symbols after the initial 9 are Reed-Solomon code symbols. Each R-S symbol can recover one known erasure; every two R-S symbols can detect and correct one other erroneous symbol. If any R-S symbols remains unused in excess of all erasures and errors, then the entire sequence can be confirmed as an R-S “codeword”, and its validity is assured, to a probability of:

$$P(1-1/2^{(5*\text{excess})})$$

For example, with one R-S symbol remaining, the probability that the EZCOD is correct is:

$$P(1-1/2^5) == .969$$

If two excess R-S symbols exist, then the probability rises to:

$$P(1-1/2^{10}) == P(1-1/1024) == 0.999$$

With 3, it's:

$$P(1-1/2^{15}) == P(1-1/32768) == 0.99997$$

Therefore, if extremely robust positions are required, select an EZCOD with 3 parity symbols, yielding almost 5-nines reliability in transmitting accurate position information – even if it must be written down, recited or entered by a human.

3.4 Precision

To identify the location of something within +/- 10 feet (3m) is simple: you must specify the Latitude (-90,90) to within 1 part in 4,194,304 (2^{22}) and Longitude (-180,180) to within 1 part in 8,388,608 (2^{23}).

The default 10-symbol EZCOD transmits 22 bits of Latitude and 23 bits of Longitude in 9 symbols of position data (the 10th is a parity symbol). The EZCOD API can encode up to 12 symbols of position data (29 bits of Latitude, and 31 bits of Longitude), yielding a maximum precision capability of +/- 20 millimeters.

Since the earth's circumference at the equator is ~40,075,000m, each part in both vertical and horizontal directions is $40,075,000 / 8,388,608 == 4.777\text{m}$. If you can specify a rectangle having sides of length equal to one part in the vertical and horizontal direction, then at the equator, you have a

square that is 4.777m on a side. So, if we know which square some geographical coordinate lies within, it is at most $\sqrt{2 * (4.777/2)^2} == 3.378\text{m}$ distant from the center of the square.

As you travel north or south, the circumference of the Longitude lines decreases, as absolute Latitude increases. The average radius of the earth is ~6,371,000m. At 53 degrees North, the circumference of the earth along a line of fixed Latitude is:

```
2 * pi * radius * cos( Latitude )
2 * 3.1415926534 * 6,371,000m * 0.60181502315
24,090,760m
```

Thus, each part along the vertical axis is still 4.777m, but each horizontal part is:

```
24,090,760 / 8,388,608 == 2.872m.
```

Now the point within each rectangle is at most:

```
sqrt( (4.777/2)^2 + (2.872/2)^2 ) == 2.787m
```

distant from the center of the rectangle.

Thus, with 9 symbols of position data, the precision of such a Latitude/Longitude encoding is at worst +/- 3.378m at the equator, at best +/-2.389m at the poles, and has an average error of less than +/-3m.

3.5 EZCOD Demo: <http://ezcod.com>

To see EZCOD in action, visit ezcod.com. Try entering:

```
R3U 1JU QUY.0
```

to see where my daughter Amarissa was born.

You can also access the Console (right click, select Inspect Element, click on "Console"), and enter the above `rskey_...` API example code.

4 RSPWD: Password Correction API

Javascript implementation of Reed-Solomon codec based password error detection and correction.

4.1 Javascript Library: js/ezpwd/rspwd.js